





10 FREE LLM RESOURCES



**From Beginner
to Advanced**



1

Large Language Model (LLM)

MacBook 



[hr Talk] Intro to Large Language Models

Intro to LLMs : Andrej Karpathy (YouTube)



2



iNeuron

Hey There, Welcome!

We are iNeuron Technology

“

O, what a tangled web we weave when first
we practice to deceive!

Walter Scott

Foundational Generative
AI : iNeuron



3



GenAI for Beginners :
Microsoft



4

FILES

Large Language Models

Text Representation

Text Generation

Deployment

Semantic Search

Prompt Engineering


Retrieval-Augmented
Generation (RAG)

Model Use

Deploy on AWS

MODULE 1.

Large Language Models



CHAPTER 1

What Are Word and Sentence
Embeddings?


[READ FULL ARTICLE ->](#)



CHAPTER 2

What is Similarity Between
Sentences?


[READ FULL ARTICLE ->](#)



CHAPTER 3

What Is Attention in Language
Models?

[READ FULL ARTICLE ->](#)



CHAPTER 4

What Are Transformer Models
and How Do They Work?

LLM University : Cohere



5

course ▾

documentation 3.8K

3,771

TRANSFORMER MODELS

Introduction

Language Processing and Language Models

Transformers, what can they do?

How do Transformers work?

What tasks do Transformers solve?

Transformer Architectures

Tokenizer

Working with LLMs

Limitations

Why

Conclusion exam

TRANSFORMERS

What to expect?

Here is a brief overview of the course:

Introduction

- Transformer models
- Using 🤗 Transformers
- Fine-tuning a pretrained model
- Sharing models and tokenizers

Diving in

- The 🤗 Datasets library
- The 🤗 Tokenizers library
- Main NLP tasks
- How to ask for help

Advanced

- Building and sharing demos
- Transformers can hear
- Transformers can see
- Optimizing for production

- Chapters 1 to 4 provide an introduction to the main concepts of the 🤗 Transformers library. By the end of this part of the course, you will be familiar with how Transformer models work and will know how to use a model from the [Hugging Face Hub](#), fine-tune it on a dataset, and share your results on the Hub!
- Chapters 5 to 8 teach the basics of 🤗 Datasets and 🤗 Tokenizers before diving into classic NLP tasks and LLM techniques. By the end of this part, you will be able to tackle the most common language processing challenges by yourself.

NLP Course : Hugging Face



6



757 views

playlist

🔖 🔄 ⋮

1



Announcing NLP Live community Sessions

Krish Naik • 222k views • 3 years ago

2



Prerequisites For NLP Live Community Session For ML And DL

Krish Naik • 68k views • 3 years ago

3



Live Day 1- Introduction And Roadmap To Natural Language Processing And Quiz

Krish Naik • 216k views • Streamed 3 years ago


4



Live Day 2- Bag Of Words, TF-IDF, Word2Vec NLP And Quiz-5000Inr Give Away

Krish Naik • 149k views • Streamed 3 years ago

5




Live Day 3- TF-IDF With Practical Application NLP And Quiz-5000Inr Give Away

Krish Naik • 87k views • Streamed 3 years ago

Natural Language Processing Playlist : Krish Naik



7



The screenshot shows the DeepLearning.AI website interface. At the top, the navigation bar includes the logo, 'Explore Courses', 'AI Newsletter' with a dropdown arrow, 'AI Dev x SF 26' with a star icon, 'Community' with a dropdown arrow, and 'Membership'. Below this, a breadcrumb trail reads 'Courses > Generative AI with LLMs'. To the right of the breadcrumb are links for 'Overview', 'Instructors' (highlighted in red), and 'FAQ'. The main heading is 'What you'll do in Generative AI with LLMs'. Below the heading is a bulleted list of course topics. At the bottom left, a small purple box contains the URL 'www.deeplearning.ai'. At the bottom right, a small grey box contains the letter 'P'.

DeepLearning.AI

Explore Courses AI Newsletter ▾ ✨ AI Dev x SF 26 Community ▾ Membership

Courses > Generative AI with LLMs Overview **Instructors** FAQ

What you'll do in Generative AI with LLMs

- Deeply understand generative AI, describing the key steps in a typical LLM-based generative AI lifecycle, from data gathering and model selection, to performance evaluation and deployment
- Describe in detail the transformer architecture that powers LLMs, how they're trained, and how fine-tuning enables LLMs to be adapted to a variety of specific use cases
- Use empirical scaling laws to optimize the model's objective function across dataset size, compute budget, and inference requirements
- Apply state-of-the-art training, tuning, inference, tools, and deployment methods to maximize the performance of models within the specific constraints of your project
- Discuss the challenges and opportunities that generative AI creates for businesses after hearing stories from industry researchers and practitioners

www.deeplearning.ai

P

GenAI with LLMs :
DeepLearning.AI



Provided proper attribution is provided, Google hereby grants permission to reproduce the tables and figures in this paper solely for use in journalistic or scholarly works.

Attention Is All You Need

Ashish Vaswani*
Google Brain
avaswani@google.com

Noam Shazeer*
Google Brain
noam@google.com

Niki Parmar*
Google Research
nikip@google.com

Jakob Uszkoreit*
Google Research
usz@google.com

Llion Jones*
Google Research
llion@google.com

Aidan N. Gomez* †
University of Toronto
aidan@cs.toronto.edu

Łukasz Kaiser*
Google Brain
lukaszkaizer@google.com

Illia Polosukhin* ‡
illia.polosukhin@gmail.com

Abstract

The dominant sequence transduction models are based on complex recurrent or convolutional neural networks that include an encoder and a decoder. The best performing models also connect the encoder and decoder through an attention mechanism. We propose a new simple network architecture, the Transformer, based solely on attention mechanisms, dispensing with recurrence and convolutions entirely. Experiments on two machine translation tasks show these models to be superior in quality while being more parallelizable and requiring significantly

06.03762v7 [cs.CL] 2 Aug 2023

Attention Is All You Need



9

LLM PowerHouse

[TWITTER](#)[LINKEDIN](#)[INSTAGRAM](#)[WEBSITE](#)

Unleash LLMs' potential through curated tutorials, best practices, and ready-to-use code for custom training and inferencing.

Overview

Welcome to LLM-PowerHouse, your ultimate resource for unleashing the full potential of Large Language Models (LLMs) with custom training and inferencing. This GitHub repository is a comprehensive and curated guide designed to empower developers, researchers, and enthusiasts to harness the true capabilities of LLMs and build intelligent applications that push the boundaries of natural language understanding.

Table of contents

- [Foundations of LLMs](#)
- [Unlock the Art of LLM Science](#)
- [Building Production-Ready LLM Applications](#)
- [In-Depth Articles](#)

LLM-PowerHouse



10



Hugging Face



Search models, datasets, users...

Trending Papers

by  AK and the research community

Papers with Code

